

Histogram Sort with Sampling

Course: COMP 5704
Parallel Algorithms and Application in Data Science
Megha Agarwal

Introduction

- What is Sorting?
- Why do we need it?
- Types of Sorting:
 - Sequential
 - Parallel
- Chief goal of parallel sorting

Input sequence

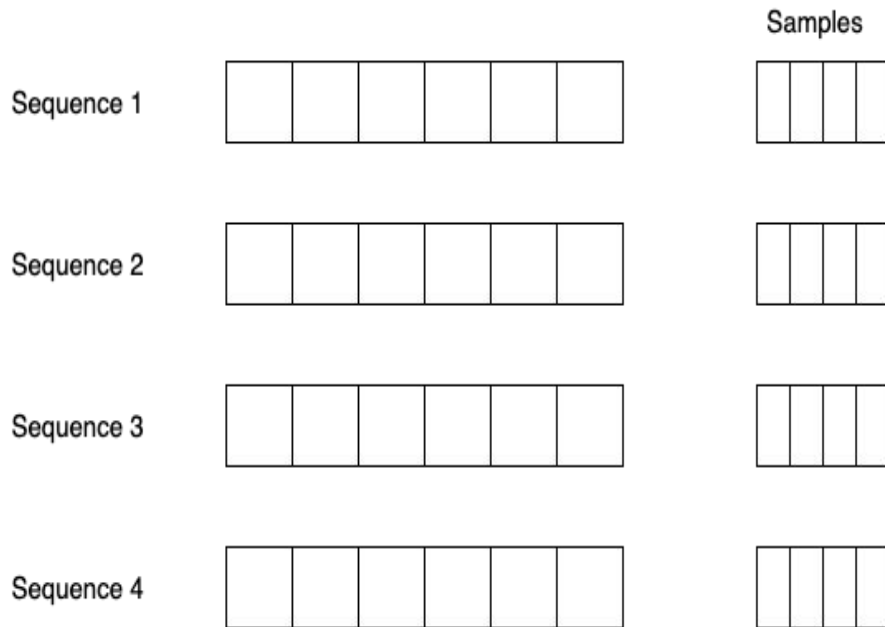
6	2	1	9	4	7	3	8	5
---	---	---	---	---	---	---	---	---

Output Sequence

1	2	3	4	5	6	7	8	9
---	---	---	---	---	---	---	---	---


Sample Sort

1. Sample: ps keys
2. Split: $p-1$ keys are selected
3. Exchange data: p^{th} bucket to p^{th} processor



Histogram Sort

1. Broadcast a probe
2. Create local histogram
3. Sum to form global histogram
4. Finalize splitters

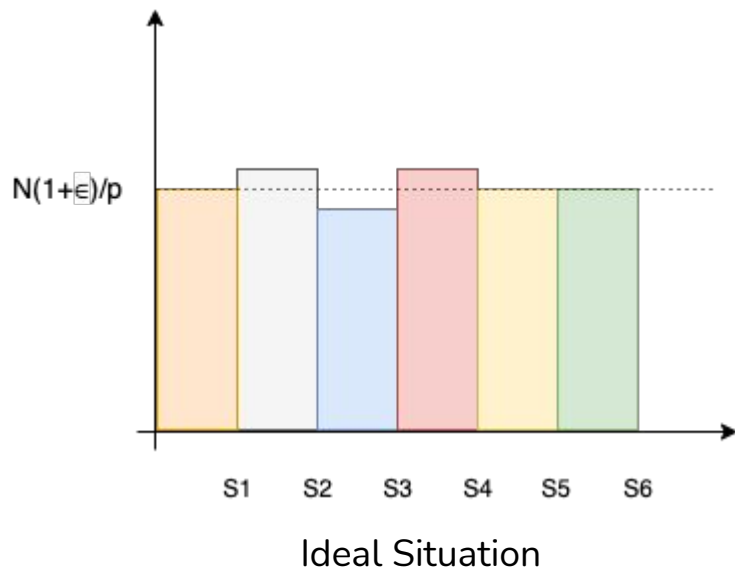



Histogram Sort with Sampling

- Goal: Approximate Splitting $\Rightarrow N(1+\epsilon)/p$
- Processor i owns all keys greater than equal to $S(i)$ and less than $S(i+1)$
:= Global balancing
- Two Major steps:
 - Histogramming
 - Sampling
- Input keys: $A(0), \dots, A(N-1)$
- Redistributed to: $I(0), \dots, I(N-1)$, where if $A(j)=I(k)$, it has rank k .
- Satisfactory splitter: $S(i) = I(\mathbf{x}(i))$, where $\mathbf{x}(i) \in \mathcal{T}_i = \left[\frac{Ni}{p} - \frac{N\epsilon}{2p}, \frac{Ni}{p} + \frac{N\epsilon}{2p} \right]$

Histogram Sort with Sampling

1. Each processor picks samples with probability ps_1/N and broadcast.
2. Create a local histogram at each processor and summed at central processor
3. Maintain a lower and upper bound $L_j(i)$ and $U_j(i)$ and update splitter intervals
4. Sample using new intervals for $j+1^{\text{th}}$ round.
5. If $j=k$,
 - a. Histogramming phase is complete and move to next step.
 - b. Else, if $j < k$, samples are collected at central processor, and move towards next round of histogramming.
6. Finally, the key closest to Ni/p is chosen as the i^{th} splitter.





Experimental setup

Charm++

- C++ based
- Supports MPI communication protocol
- Divides into processor elements called *chares*.
- Steps:
 - Local Sorting
 - Splitter Determination
 - Data exchange

Expected results of HSS

- Sampling ratio $s = O\left(p \sqrt[k]{\frac{\log p}{\epsilon}}\right)$ and $k = \log(\log p / \epsilon)$
- $O(p)$ samples can achieve global sorting in $O(\log N/p + \log \log p)$ rounds
- Costs:
 - Computation in local Sorting: $O((N/p) \log N/p)$
 - Sampling: $O(S)$ at local processor and $O(S \log p)$ for sorting at central processor
 - Computing local histogram: $O(S \log N/p)$
 - Computation of sampling and histogramming per stage: $O(r \log((\log r)/\epsilon)) \log N$
- Communication overhead due to multiple stage sorting.

HSS compared to other algorithms

AMS-sort

- Better for one round of histogramming by $\Theta(\min(\log p, 1/\epsilon))$
- HSS achieves a globally-balanced splitting, making it easily generalizable
- Takes approximately 3x time for splitting phase than HSS

HykSort

- Requires at least $\Omega(\log(p)/\log^2 \log(p))$ times more samples.
- Faster convergence of splitters

Thank you!